



# SCALEOUT

Configuration Management Tools -  
Software-Defined Networks



Max Andersson

- Primary goal of SDN
  - Open & Programmable
    - Traffic Engineering
    - Security
    - Qos
    - Routing
    - Switching
    - Virtualization
    - Monitoring
    - Load Balancing
    - New applications

- 3 Layer Model OS
  - South side
  - North side

- 3 Layer SDN model
  - Forwarding Device Southside
  - Net applications
  - NEtwork operating system
- Fast path , caches paths. Can ask operation system

- Forwarding devices / Data Plane
  - Can be hardware switches if they support openflow etc, but can also be software devices.
- Southbound Interface
  - OpenFlow
  - OVSDB -(Open vSwitch)
  - NETCONF
  - SNMP

- Application interfaces
  - Java Api
  - Northbound (e.g RESTConf)
- Layer details
  - Slicing, forwarding to different network operating systems
  - Logically centralized vs physically centralized
  - Clustered Network operating system.
  - Regions with EAST/WEST protocol
  - Hierarchies

- Availability and scalability
- SDN vs Traditional networks
  - Traditional networks has a control plane and data plane in one physical system
- BGP, Multiprotocol Label Switching (MPLS), Open Shortest Path First (OSPF)
  - Control plane pushes policies down to the fast path of the data plane.

- Openstack
  - Networks
  - Segmentation Methods
    - VLAN's
    - VXLAN
    - GRE
    - Network Namespaces
    - OpenFlow Rules
    - Subnet
      - Default DHCP service (dnsmasq)'



- **Kubernetes**
  - Flannel
    - Why
      - Layer 2 solution
      - Simple & Mature
      - Overlays are useful when network address space is limited
      - Overlays auto-configure
    - Scenario
      - On-Prem or custom cloud where native routing isn't possible
    - Why not
      - Native routing is faster and easier to debug.
      - You need Calico if you want network policies.

## Calico

- **Why**
  - Layer 3 solution
  - Good Network policy support
  - Default on most kubernetes distributions
  - Easy to debug on hosts by looking on route tables. BGP allows access inside and outside the cluster.
- **Scenario**
  - On-Prem with native routing or cloud kubernetes services
- **Why Not**
  - IP-in-IP mode is needed when routing between subnets, which negates some of the performance benefits vs an overlay
  - BGP can be scary.

## Calico

- **Why**
  - Layer 3 solution
  - BGP
  - Good Network policy support
  - Default on most kubernetes distributions
  - Easy to debug on hosts by looking on route tables. BGP allows access inside and outside the cluster.
- **Scenario**
  - On-Prem with native routing or cloud kubernetes services
- **Why Not**
  - IP-in-IP mode is needed when routing between subnets, which negates some of the performance benefits vs an overlay
  - BGP can be scary.

- **Kube-Router**
  - **Why**
    - **Layer 3 solution**
    - **BGP**
    - **Single Go binary built from the ground up for Kubernetes.**
    - **Uses new IPVS/LVS kernel features to improve service load balancing performance.**
    - **Also does direct server return to improve latency.**
  - **Scenario**
    - **On-Prem or custom cloud latency focus.**
  - **Why not**
    - **Similar to Calico in that it uses IP-in-IP by default to encapsulate traffic between subnets.**
    - **Quite a new project and although it's in use in production at some companies it's still not v1.**

**Net-  
App**

**Net-  
App**

**Net-  
App**

**Net-  
App**

**Net-  
App**

**Network Operating System**

**Forwarding  
device**

**Forwarding  
device**

**Forwarding  
device**