

SCALEOUT

Configuration Management Tools -
Observability



Max Andersson

- The Need for observability
 - Not purely an operational concern.
 - Observability is not only about logs metrics and traces.

- Observability is a property of a system that has been
 - Designed, build, tested, deployed, operated, monitored, maintained and evolved.

- Observability because
 - No complex system is ever fully healthy
 - Distributed systems are unpredictable
 - Impossible to predict all states of partial failures the system can end up in.
 - Failure needs to be embraced at every phase, from design to operation.
 - Ease of debugging is a important for maintenance of an evolving system.

- The Three Pillars of Observability
 - Logs
 - Metrics (Telemetry)
 - Traces

- Monitoring and Observability
 - Alerting based on monitoring data
 - Blackbox vs Whitebox monitoring
 - Monitoring should provide an overview of health.
 - Ability to drill down on component level
 - Monitoring is a part of the Security process

- Alerting
 - Event Failures should immediately provide visibility of impact.
 - All alerts and signals that generate them should be actionable.

- Failure modes
 - Tolerated
 - Alleviated
 - Graceful degradation
 - Backpressure
 - Retries
 - Timeouts
 - Circuit braking
 - Rate limiting
 - Triggered

- Monitoring Signals
 - USE Metrics
 - Free Memory (**U**tilization)
 - CPU queue length (**S**aturation)
 - Device errors (**E**rrors)
 - RED Metrics
 - Request rate
 - Error rate
 - Duration of requests
 - Golden Signals (minimum viable signals)
 - Latency
 - Errors
 - Traffic
 - Saturation

- Coding for failure
 - Semantics
 - Characteristics
 - Debuggable

- Coding for failure
 - Semantics
 - Characteristics
 - Debuggable

- Semantics
 - How a service is deployed and what tools
 - Is the service binding to a standard port ?
 - How an application handles signals
 - How processes start on a given host
 - How it registers with service discovery
 - How graceful (or not) the restarts are.
 - How configuration is fed to the process
 - Concurrency model of the application
 - Multithreaded, single threaded, event driven, actor based. Etc.
 - How the Reverse proxy/ Ingress handles connections

- What are logs?
 - Historical Records of events.
 - Records events and status of systems in a time sequential format
 - Records Activity on a system/network
 - Provides an Audit trail of who, what, where, when and why.

- Why are they important?
 - Determining what happened - Audit trail
 - Intrusion Detection
 - Incident Containment
 - Forensic Analysis
 - Proactive Analysis
 - Real-time alerts
 - Determining the health of a network
 - Debugging/Trubbleshooting
 - Proactive maintenance

- Log Types
 - Event Logs
 - Immutable
 - Timestamped records
 - Discrete Events
 - Formats
 - Plaintext
 - Structured
 - Binary
 - Replication
 - Point-in-time recovery
 - Systemd journal logs
 - Pflog

- Type of Logs
 - Traffic Logs
 - Threat Logs
 - URL Filtering Logs
 - Data Filtering Logs
 - Correlation logs
 - Tunnel inspection logs
 - Config logs
 - System logs
 - User Logs
 - Alarm Logs
 - Authentication logs

- Where to find the logs?
 - Operating Systems
 - Applications
 - Device Logs
 - Routers
 - Firewalls
 - Switches
 - etc..

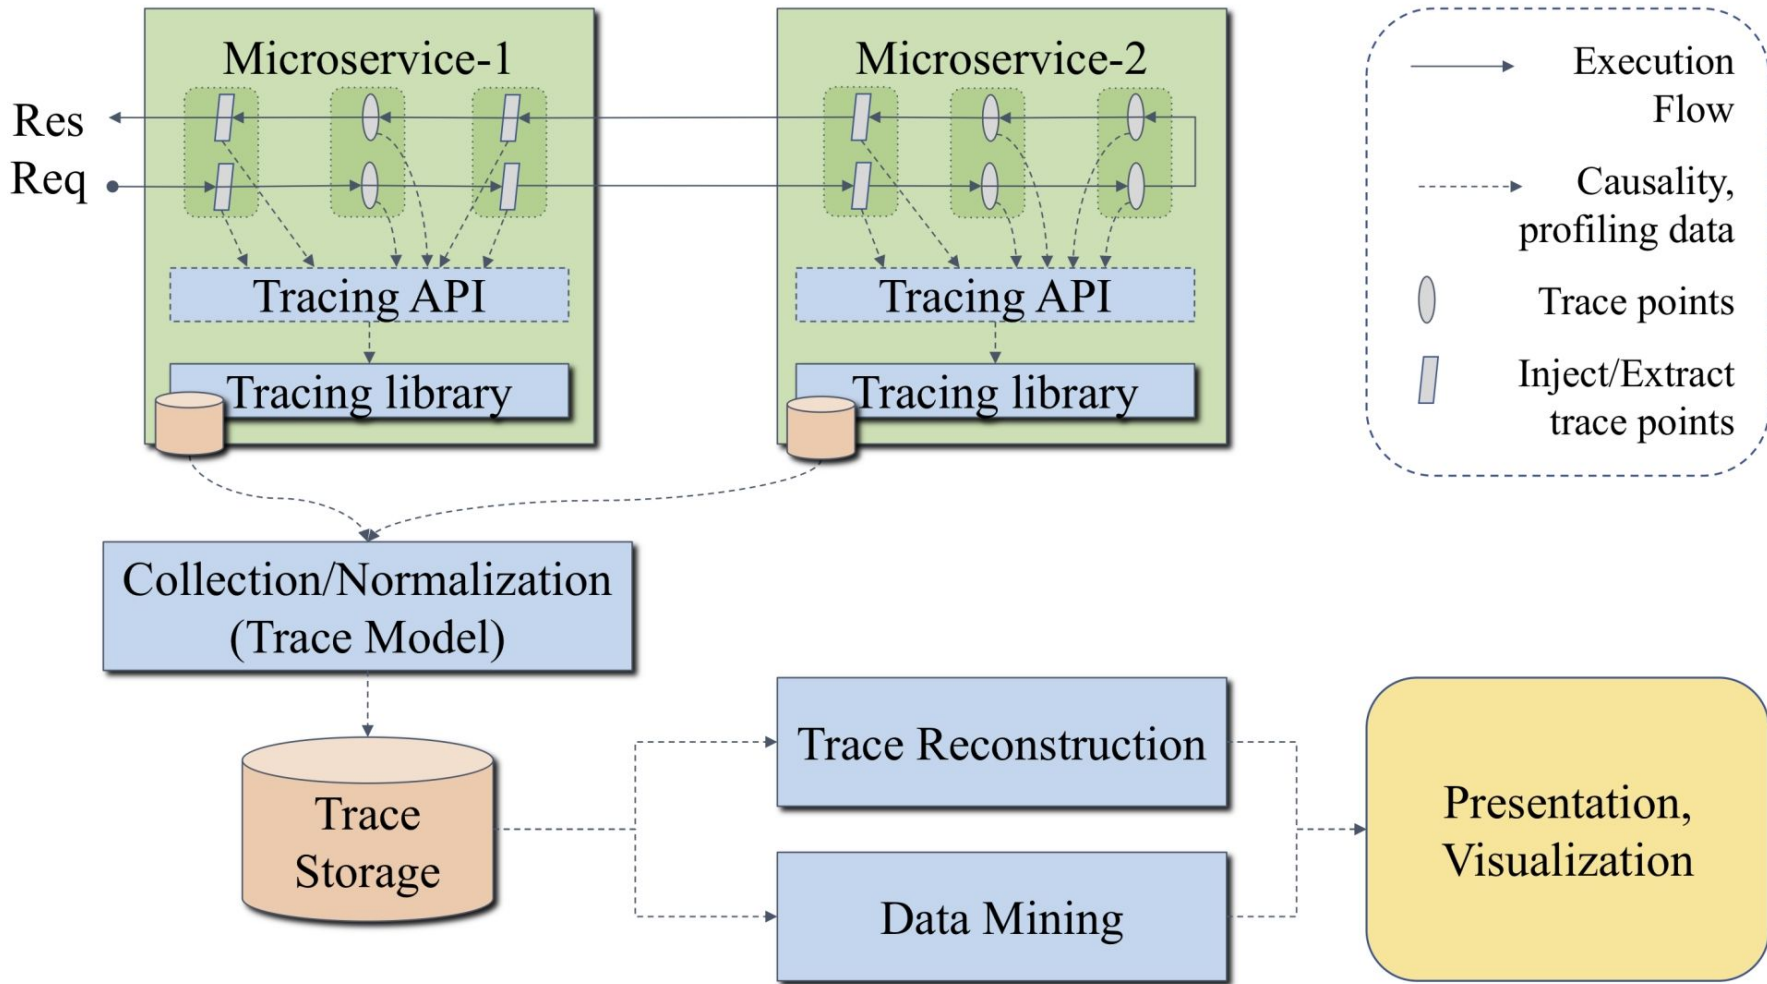
- Challenges
 - Different log formats
 - Logs are written by developers
 - Format and content can be hard to obscure and hard to understand.
 - Information Volume
 - Anomalies Detection is hard!

- Best Practices
 - Centralized log collection.
 - Secure transmission of logs
 - Normalize the data to same format
 - Review the logs
 - Manually and Automatically
 - Log Rotation
 - Time schedule
 - Naming Convention
 - Log Retention
 - Disk Space
 - Regulatory requirements?
 - Archive

- So....
 - What to log? **Logging Policy**
 - What to centralize? **Log collection**
 - What to save? **Log retention**
 - What to look at? **Log Review**
 - What to alert on? **Log Monitoring**


- Tracing
 - A Trace is a representation of a services of causally related distributed events that represents the end-to-end request flow.
 - Identify a specific point in request flow.
 - Zipkin & Jaeger are the most popular OpenTracing-compliant open source distributed tracing solutions

- Distributed Tracing
 - Tracing Infrastructure attaches Contextual metadata to each requests.
 - Trace points in code, records events to be annotated.
 - Events are tagged with context and causality references generated by previous events




- Service meshes
 - Makes implementing tracing easier
 - The data plane implements the tracing and collection of statistics from a proxy.
 - Getting support of all of its services can mean requirements of the application layer


No Concurrency




Basic Concurrency



Async Concurrency



Distributed Concurrency



Observability